


Does the model reason more when it believes it is talking to the CEO?

(gap is positive and clear of zero on all five models; label at each point is the mechanism verdict — see legend)



R1 = CEO minus analyst reasoning tokens per sample.

"Models that motivated the test" first suggested this effect; "held-out models" tested it on fresh models afterwards.

"more deliberation" = the gap holds up even after controlling for longer multi-turn episodes.

"verbosity, not separable" = the model may simply be writing more, not thinking more.