


Does the model reason more when it believes it is talking to the CEO?

(gap is positive and clear of zero on all five arms; label at each point is the mechanism verdict — see legend)



Reasoning-token gap: CEO minus analyst, as % of the analyst's mean (labelled R1 in the underlying analysis)
95% scenario-clustered bootstrap CI (7 clusters, 10,000 draws, seed 6)

R1 = CEO minus analyst reasoning tokens per sample.

"Arms that motivated the test" first suggested this effect; "held-out arms" tested it on fresh models afterwards.

"more deliberation" = the gap holds up even after controlling for longer multi-turn episodes.

"verbosity, not separable" = the model may simply be writing more, not thinking more.