


CEO minus analyst, harmful-action rate, by scenario — descriptive only
(status-irrelevant scenarios; n = 20 per cell on the five main arms; no intervals shown)



"no harm": the arm produced no harmful episode in that scenario, so there was no opportunity to defer. "not scored": the harmful outcome cannot be decided from the recorded tool arguments for this scenario on the earlier harness. gpt-5-nano (grey label) ran on an earlier harness version with n = 9-10 per cell, not 20, using a backfilled harm outcome, and is not a sixth directly comparable arm.