

Figure 3. CEO minus analyst, harmful-action outcome, one scenario at a time
status_irrelevant only · n = 20 per cell on the five clean arms (nano cells state their own n) · ILLUSTRATIVE, no intervals



"no harm": the arm produced no harmful episode in that scenario, so there was no opportunity to defer. "n/a": structurally undecidable under AI-23. gpt-5-nano (grey label) is a legacy cross-fork arm, n=9-10 per cell not 20, with a backfilled harm outcome, and is not a sixth clean point.